

AI Illustrator

The machine that turns a manuscript into a finished picture book



In twelve weeks you receive both books finished and print-ready, and you own the tool that produced them. From that point, adding a further title becomes a matter of days rather than months.

DELIVERED

Two finished books

ELLA AND BILBO, PRINT-READY

DELIVERED

The AI illustrator

REUSABLE ON EVERY FUTURE TITLE

What we are actually building

You receive both books finished and print-ready, and the tool that produced them. It takes a manuscript in at one end and hands a print-ready picture book out at the other.

It reads the manuscript

The story is broken into pages, and the illustration notes are treated as binding instructions rather than loose suggestions.

It keeps a fixed cast

Each character is drawn once, held to exact colours and clothing, and reused, so nobody changes shape between pages.

It composes each page

Backgrounds and characters are assembled with a quiet area reserved for the words before anything is drawn.

It sets the words

Captions and bubbles are placed where they can be read, checked for contrast, and kept word for word to the manuscript.

It builds the book

Front matter, story, back matter and cover are assembled at one trim size and exported ready for the printer.

It marks its own work

An independent scorecard reads the finished book page by page and reports what still falls short.

What the tool is worth after delivery

Repeatable, not one-off.

The same steps run again on a manuscript nobody has seen yet.

Built in replaceable parts.

A better method for one step drops in without rebuilding the rest.

The value compounds.

The books hold their value as titles, and the tool adds to it on every commission after.

It keeps your decisions.

Cast, colours and layout rules stay stored, so the next book starts where this one finished.

It checks itself.

The scorecard measures a new book before anyone reads it.

Days, not months.

The next title starts with the cast built and the rules already written.

Where the AI illustrator stands today

You cannot improve a machine you are not measuring. These scores come from a page-by-page review of 39 pages it has produced, read by eye against the character reference sheets. They are deliberately not the machine's opinion of itself, because that self-check only looks at what it was told to look at.



A note on the numbers

These targets are set to be realistic rather than flattering. Every measure is scored out of 100, and 90 is the level we believe is genuinely achievable in twelve weeks, so it is the figure we are willing to be held to. Each measure reaches 90 in the week whose work owns it and then holds there. Book structure is the exception. It is a pass or fail check on every page, and it has to come back clean before delivery.

The six faults in it today

Each is a gap in the machine, not a mistake on one page, so fixing it once fixes it for every book that follows. Each fault is owned by a single week, in the order a reader would notice it.

01

Characters change between pages

Size, face and clothing shift from page to page. Ella appears twice on one page, Obi loses his hat, and Mom's red jersey comes and goes.

ADDRESSED IN WEEKS 2 AND 3

02

Words land on busy artwork

Text is placed over fences, crowds and shopfronts because no quiet area was planned for it before the picture was drawn.

ADDRESSED IN WEEK 4

03

Some captions are hard to read

One caption is scrambled into nonsense, and white text over a pale sky has too little contrast to read comfortably.

ADDRESSED IN WEEK 5

04

Working notes leak onto the page

Internal instructions have been printed as if they were part of the story, and the back cover is still a placeholder.

ADDRESSED IN WEEK 6

05

No speech or thought bubbles

Dialogue is only ever narrated. The bubbles readers expect in this kind of book were never built.

ADDRESSED IN WEEK 7

06

Scenes and props miss the notes

Invented extras appear that were never asked for, such as a hot dog stand, and some figures are stretched, duplicated or cropped off awkwardly.

ADDRESSED IN WEEK 8



The one change that unlocks the rest

Drawing each character separately and placing them into the scene fixes their size and appearance, and records where each one sits, which is what makes quiet space and readable text possible.

WEEKS 2 TO 5 BUILD ON THIS




Why the build weeks are not redrawn twice

Weeks 2 to 8 each retire one fault, applied only to the pages it touches. Week 9 runs the finished process end to end once, so no book is ever redrawn twice.

ONE FAULT PER WEEK, IN ORDER

How the twelve weeks flow

Every stage upgrades the AI illustrator, and nothing later depends on a promise, only on work that has already passed its check.



Week by week

Each week has one headline goal and one outcome, and does not close until that outcome is met or the miss is written down.

WEEK ONE · GROUNDWORK

Measure honestly before building anything

- Turn the six problems above into a repeatable scorecard that reads a finished book page by page.
- Keep the scorecard separate from the software that makes the art, so it is free to disagree with it.
- Add a short human review sheet for the things software judges poorly, such as readability and taste.
- Save today's books as the frozen starting point that all progress is measured against.

OUTCOME

The scorecard reproduces the hand review within five points on both books, and the report is saved to file rather than shown on screen and lost.

WEEK TWO · BUILDING

Run both full books through the new drawing method

- Draw each main character once in every pose the books need, then reuse those drawings everywhere.
- Draw each background without characters in it, then place the characters into the scene at a set size and position.
- Record exactly where each character sits on each page, which every later week relies on.
- Make this the standard way both books are produced from now on.

OUTCOME

Both books are rebuilt on the new drawing method, character size varies by no more than five percent across pages, and the duplicate Ella and missing hat problems are gone.

WEEK THREE · BUILDING

Hold the people and the look-alikes steady

- Lock every recurring person to exact colours and clothing, so an outfit cannot drift between pages.
- Bring the people who appear on three or more pages into the same reuse method as the animals.
- Create a paired reference for characters who look similar, so the two dogs and the two girls stay distinct.
- Extend the scorecard to check clothing specifically.

OUTCOME

Character consistency reaches 90 out of 100, with no clothing drift on any named character. This measure now holds at 90 for the rest of the project.

WEEK FOUR · BUILDING

Design a calm area into every page

- Decide where the words will go before the picture is drawn, then ask for open sky, soft ground or plain wall in that spot.
- Use the known character positions from week 2 to keep faces and bodies out of the text area.
- Raise or lower the horizon where it helps create that open space naturally.
- Check how busy the text area actually is, and redraw the page if it is too cluttered.

OUTCOME

Calm space reaches 90 out of 100, every page passes the clutter check in its text area, and the measure holds at 90 from here on.

WEEK FIVE · BUILDING

Rebuild how the words sit on the art

- Fix the fault that scrambled a caption into nonsense.
- Place text using the known character positions and the calm area, at one consistent size across the book.
- Guarantee readable contrast by adapting the ink colour and adding a soft glow behind it, with no boxed label look.
- Carry the manuscript wording through unchanged, so what is printed matches the author's text exactly.

OUTCOME

Readability reaches 90 out of 100, contrast passes on every page, no caption is scrambled, cut off or split across the spine, and the measure holds at 90.

WEEK SIX · YOUR FIRST LOOK

Fix the book structure and send you one complete book

- Write the real front and back matter, including title page, dedication, about the author, a note on fostering and volunteering, and the back cover.
- Add a hard stop that fails the build if any internal note reaches a printed page.
- Assemble the book in manuscript order at one consistent trim size.
- Send you the stronger of the two books in full, so direction can be confirmed at the halfway mark rather than at the end.

OUTCOME

Structure comes back clean with no leaked notes and correct front and back matter, and your first round of feedback is logged.

WEEK SEVEN · BUILDING

Add speech and thought bubbles

- Separate spoken lines from narration and from a character's private thoughts.
- Draw speech bubbles with a tail pointing at the speaker, and thought bubbles as clouds with trailing dots.
- Keep bubbles clear of faces, key artwork and the calm text area.
- Check that each bubble is readable and points at the right character.

OUTCOME

Bubbles reach 90 out of 100 wherever dialogue exists, hold at 90, and can be switched off per book if you prefer them used sparingly.

WEEK EIGHT · BUILDING

Make every page match the location and props

- Treat the illustration notes in the manuscript as binding, so the location and named props are not open to interpretation.
- Reject invented extras that were never asked for, such as a hot dog stand that does not belong in the scene.
- Catch stretched limbs, duplicated figures and subjects cropped off awkwardly.
- Keep faces and key objects away from the centre fold on double page spreads.

OUTCOME

Scene accuracy reaches 90 out of 100, every page matches the location and props its notes call for, with no warped, duplicated or badly cropped artwork.

WEEK NINE · BUILDING

Run both books end to end on the finished process

- Produce both books in a single integrated pass on the finished process, so the parts are proven to work together and not just on their own.
- Score every page on all six measures and rank any remaining gaps by how much a reader would notice them.
- Fix the worst first and rerun until every measure holds at its target.

OUTCOME

Both books hold the week 12 targets on our own scorecard, produced in one integrated run rather than a second full redraw, before you are asked to look at them.

WEEK TEN · YOUR REVIEW

Both complete books go to you

- Deliver both full books, with a short note explaining what changed since the starting point and by how much.
- Sort every comment you make into a category, so recurring themes are handled once rather than page by page.
- Agree together on what must be fixed, what would be nice to fix, and what falls outside this project.

OUTCOME

A complete, agreed and prioritised list of your feedback, with an estimate against each item.

WEEK ELEVEN · BUILDING

Work through your notes without breaking what already works

- Complete every item on the agreed must-fix list.
- Rerun the full review after each batch, because late changes have previously undone earlier wins.
- Write any new preference of yours into the standing rules, so it is applied automatically from then on.

OUTCOME

The agreed must-fix list is closed, no measure has slipped below its week 9 level, and you are happy with a preview.

WEEK TWELVE · DELIVERY

Print preparation and hand-over

- Prepare proper print files with the correct trim, bleed and safe margins, at print resolution and in print colour.
- Run the final review on the print files themselves, because correct on screen is not the same as correct on press.
- Hand over the machine itself, together with the fixed cast, the scorecard and a one page guide to running it on the next manuscript.
- Hand over the print files, the source files and the final report alongside it.

OUTCOME

The machine is handed over and proven on a new manuscript, both books meet their targets, the print specification is validated, and you sign off.

The progress you should expect to see

Scores are out of 100. Each measure climbs in the week whose task owns it, reaches 90 there, and then holds. Book structure is a pass or fail check rather than a score.

Red is still short of the target. **Teal** is the week a target is reached and held.


Week	Milestone	Characters	Calm space	Readability	Bubbles	Scene	Structure
0	Starting point, measured by hand	75	70	72	none	72	failing
2	New drawing method on both books	84	72	72	none	74	failing
3	People and clothing held steady	90	76	73	none	75	failing
4	Calm space designed in	90	90	80	none	78	failing
5	Words rebuilt on the art	90	90	90	none	82	partial
6	Structure fixed, first book to you	90	90	90	none	84	clean
7	Bubbles added	90	90	90	90	86	clean
8	Scenes and props matched	90	90	90	90	90	clean
9	Final integrated production run	90	90	90	90	90	clean
10 to 11	Your review and the fixes	90	90	90	90	90	clean
12	Print ready and handed over	90	90	90	90	90	clean

What 90 looks like on the page

At 90 a reader gets through the book without being pulled out of the story by anything on it. Past that point the remaining calls are made by eye rather than by measurement, which is why the final pass before delivery is a person reading the book rather than a score.

Where the twelve weeks actually go

Rebuilding how the artwork is made takes the largest share, because character consistency is the keystone every later week reuses.



What could slow us down

Each of these is known in advance and has a planned response.

Placed characters could look pasted on

Drawing characters separately risks a cut and paste feel on busy pages.

Response. A blending step, a size check in week 2, and the older drawing method kept available as a fallback for individual pages.

A pose we need might be missing

If a character was never drawn crouching, that page ends up awkward.

Response. Audit the pose set against both manuscripts in week 2 and draw extras on demand.

Bubbles could crowd the text

Bubbles and captions compete for the same clear space on a page.

Response. Bubbles are built in week 7, after the text placement work, so they reuse the same layout rules rather than fighting them.

Late feedback could arrive in bulk

A large set of notes at week 10 would compress the remaining schedule.

Response. The week 6 checkpoint puts a complete book in front of you at the halfway mark, so direction is confirmed early.

Late fixes could undo earlier wins

This has happened before, where a fix in one place reopened a problem elsewhere.

Response. The full review is rerun after every batch of fixes in week 11, and no measure is allowed to slip below its week 9 level.

Content filters could block a reference

Automated safety filters occasionally reject images of children or animals.

Response. The existing rewording and fallback handling stays in place, and this has been low impact so far.

What you own at the end

The tool comes first on this list on purpose, because it is the part that keeps paying. Nothing is called complete until every item holds.

- ✓ **The AI illustrator itself**, able to take a new manuscript and return a finished book with only light touch-up by hand.

- ✓ **The fixed cast**, so returning characters look the same in the next book as they did in this one.

- ✓ **The scorecard**, so the quality of any future book can be measured rather than argued about.

- ✓ **A short guide** covering how to run the machine on the next title.

- ✓ Both books scoring 90 out of 100 on all five measured dimensions.

- ✓ No internal notes on any page, with correct front and back matter and one consistent trim size.

- ✓ Every caption readable, contrast checked, clear of the centre fold and word for word to the manuscript.

- ✓ Print files at print resolution and colour, with the correct bleed and trim, signed off by you.

X

What we need from you

Three moments where your input keeps the schedule on track.

FIRST LOOK

Week Six

Review one complete book and confirm the overall direction. Roughly two to three days of your time.

FULL REVIEW

Week Ten

Review both full books and give consolidated notes. Roughly one week of your time.

SIGN-OFF

Week Twelve

Approve the print files for hand-over. Roughly one to two days of your time.

X I

What happens on the book after these two

This is the return on the twelve weeks, and the reason the machine matters more than the two books it produces.

The cast is already built

Any character who returns is drawn and locked already, so a sequel starts with its cast in place rather than from a blank page.

The rules are already written

Every preference you gave us in weeks 6 and 10 stays encoded, so the next book starts finished instead of relearning your taste.

The scorecard travels with it

A new book can be measured on the same six dimensions before anyone reads it, so quality is checked rather than hoped for.

Read this plan as the build log for the AI illustrator, with both finished books delivered alongside it. Starting point figures come from a page-by-page visual review of what the machine has produced so far, not from its own self-report. Every weekly finish line is measured by the independent scorecard built in week 1, which is deliberately able to disagree with the part of the machine that makes the artwork. Weeks 2 to 8 each retire one fault, week 9 runs the finished process once, and weeks 10 and 11 are for your review and its fixes.