

Request: GPU Fine-Tuning Access for Character Consistency

To make our picture-book characters look identical on every page, we need to fine-tune the image model per character. That needs a small pay-as-you-go GPU training service (~\$10 one-time + a few cents per image). Recommended: **fal.ai** or **Replicate** — or our own cloud/GPU if we have one.

The ask

Approval for a **pay-as-you-go account on a GPU/AI fine-tuning service** (no monthly subscription, no lock-in). First choice **fal.ai**; **Replicate.com** is an equivalent alternative. Estimated cost below.

The problem we're solving

Our tool generates illustrated children's books in which the **same characters** (two dogs, their owners, and a mascot) must appear on **~22 pages and look identical every time** — same coat colour, same cap, same size.

Every off-the-shelf image model draws each picture **from scratch**, so the same character comes out looking different on every page (wrong colour, cap appearing/disappearing, changing size/breed). We have spent considerable effort trying to solve this with detailed prompts and reference images. It has improved, but it **cannot be fully solved that way** — it is a structural limitation of how these models work.

Why a “better model” does not fix this

This is the key point:

It is not a model-quality problem — it is a model-capability problem.

- **Every** text-to-image model — DALL·E, Midjourney, Google Imagen, Flux, Stable Diffusion — shares the same limitation: each image is drawn freehand, so none of them can hold an exact character across many images on their own. A more advanced model produces **prettier** pages but is **just as inconsistent**.
- The proven, industry-standard fix is to **fine-tune the model on each character** (a lightweight add-on known as a “LoRA”). This teaches the model exactly what each character looks like, so it renders them the same on every page — while still painting them naturally into any scene or pose.
- **Fine-tuning requires GPU training infrastructure.** Our current provider (OpenRouter) only *runs* models; it cannot *train* them. That single missing capability is the entire reason for this request.

What the money buys

Item	Cost
One-time training of all 5 characters	~\$10 total (~\$2 each)
Ongoing image generation	~\$0.03–0.05 per page image
Monthly commitment	None — pay only for usage

Training each character takes ~3–5 minutes. This is a one-time setup per book; re-runs cost only pennies.

Options (any one of these works)

1. **fal.ai** (*recommended*) — fastest and cheapest for this specific task; pay-as-you-go.
2. **Replicate.com** — equivalent capability and price; direct substitute.
3. **Our existing cloud**, if we have one — AWS SageMaker, Google Vertex AI, or Azure ML can all fine-tune image models (preferable if there's already a cloud contract/credits).
4. **A company GPU / workstation** — if we have a machine with a capable NVIDIA GPU, we can train locally at no per-use cost.

Recommendation

Approve a **pay-as-you-go fal.ai (or Replicate) account**, or point me to our existing cloud/GPU if the company already has one. The total spend to prove this out is ~\$10, and it is the only reliable way to deliver the character consistency the books require.