

Why We Should Use LoRA for Character Consistency

The problem

Our tool makes illustrated children’s books where the **same characters** — the two dogs, their mom and dad, and the mascot — appear on ~22 pages and **must look identical every time**: same face, same coat, same cap, same size.

Ordinary AI image tools draw every picture **from scratch**, so the same character comes out slightly different on every page — the coat colour shifts, the cap appears or disappears, the size changes, the face drifts. We have tried hard to fix this with detailed prompts, reference pictures, and automatic corrections. It helped, but it **cannot be fully solved that way** — it is a built-in limitation of how these tools work.

Why the two methods we already tried each fall short

We built and tested both realistic in-house approaches. Each fixes one half of the problem and breaks the other:

Approach	Same character every page?	Natural-looking scenes?
Draw each page from scratch	No — the character drifts	Yes — scenes look natural
Reuse the same cut-out character	Yes — identical	No — looks pasted / stiffly placed

So we are stuck choosing between *consistent characters* **or** *natural pictures*. We cannot get both this way.

What LoRA is (in plain terms)

A **LoRA** is a small “memory add-on” for the image tool. We show it about 10–20 pictures of one character, and it **teaches the tool what that character looks like** — one small add-on per character (five in total for this book).

After that, whenever we ask for a picture, the tool **already knows** the character and paints them the same way every time, while still drawing each scene freshly and naturally.

Why LoRA gives the best result

It is the only method that delivers **both** at once:

Approach	Same character	Natural scenes
Draw from scratch	No	Yes
Reuse cut-outs	Yes	No
LoRA	Yes	Yes

The reason is simple: with LoRA the character’s identity is **learned by the tool itself**, not re-described in words (which it interprets loosely) or pasted in as a fixed picture (which it can’t re-pose). So the character is **painted fresh into every scene** — correct pose, angle, lighting, and interaction — **and** still comes out looking like the same character. This is the **industry-standard way** professionals keep a character consistent across many AI images.

What it costs and what we need

- **A pay-as-you-go account** on a fine-tuning service (**fal.ai** recommended, or Replicate) — no monthly fee, no lock-in. This is the one thing we don’t have yet; our current provider can only *run* models, not *teach* them.
- **One API key** added to our settings.
- **Training pictures** — we already have most of these from work we’ve done.

Item	Cost
Teach all five characters (one-time)	~\$10 total (~\$2 each)
Each finished page	~\$0.03–0.05
Monthly commitment	None — pay only for what we use

A full book re-render costs about **\$1–2**, and re-runs after that cost pennies.

The honest caveats

- It needs that small external service (**~\$10**) — that is the only real blocker.
- It is **very** consistent (~95%+), not a guaranteed 100%; a rare page may still need a quick redo. But it is by far the most reliable option available.
- We will re-establish the art style once, since LoRA runs on a different (equally capable) image engine.

Recommendation

Approve a small **pay-as-you-go fal.ai (or Replicate) account**. It is the only proven way to get **consistent characters *and* natural illustrations** — exactly what the books require — and the total spend to prove it out is about **\$10**. Almost everything else is already built; we just need the account and key to begin.